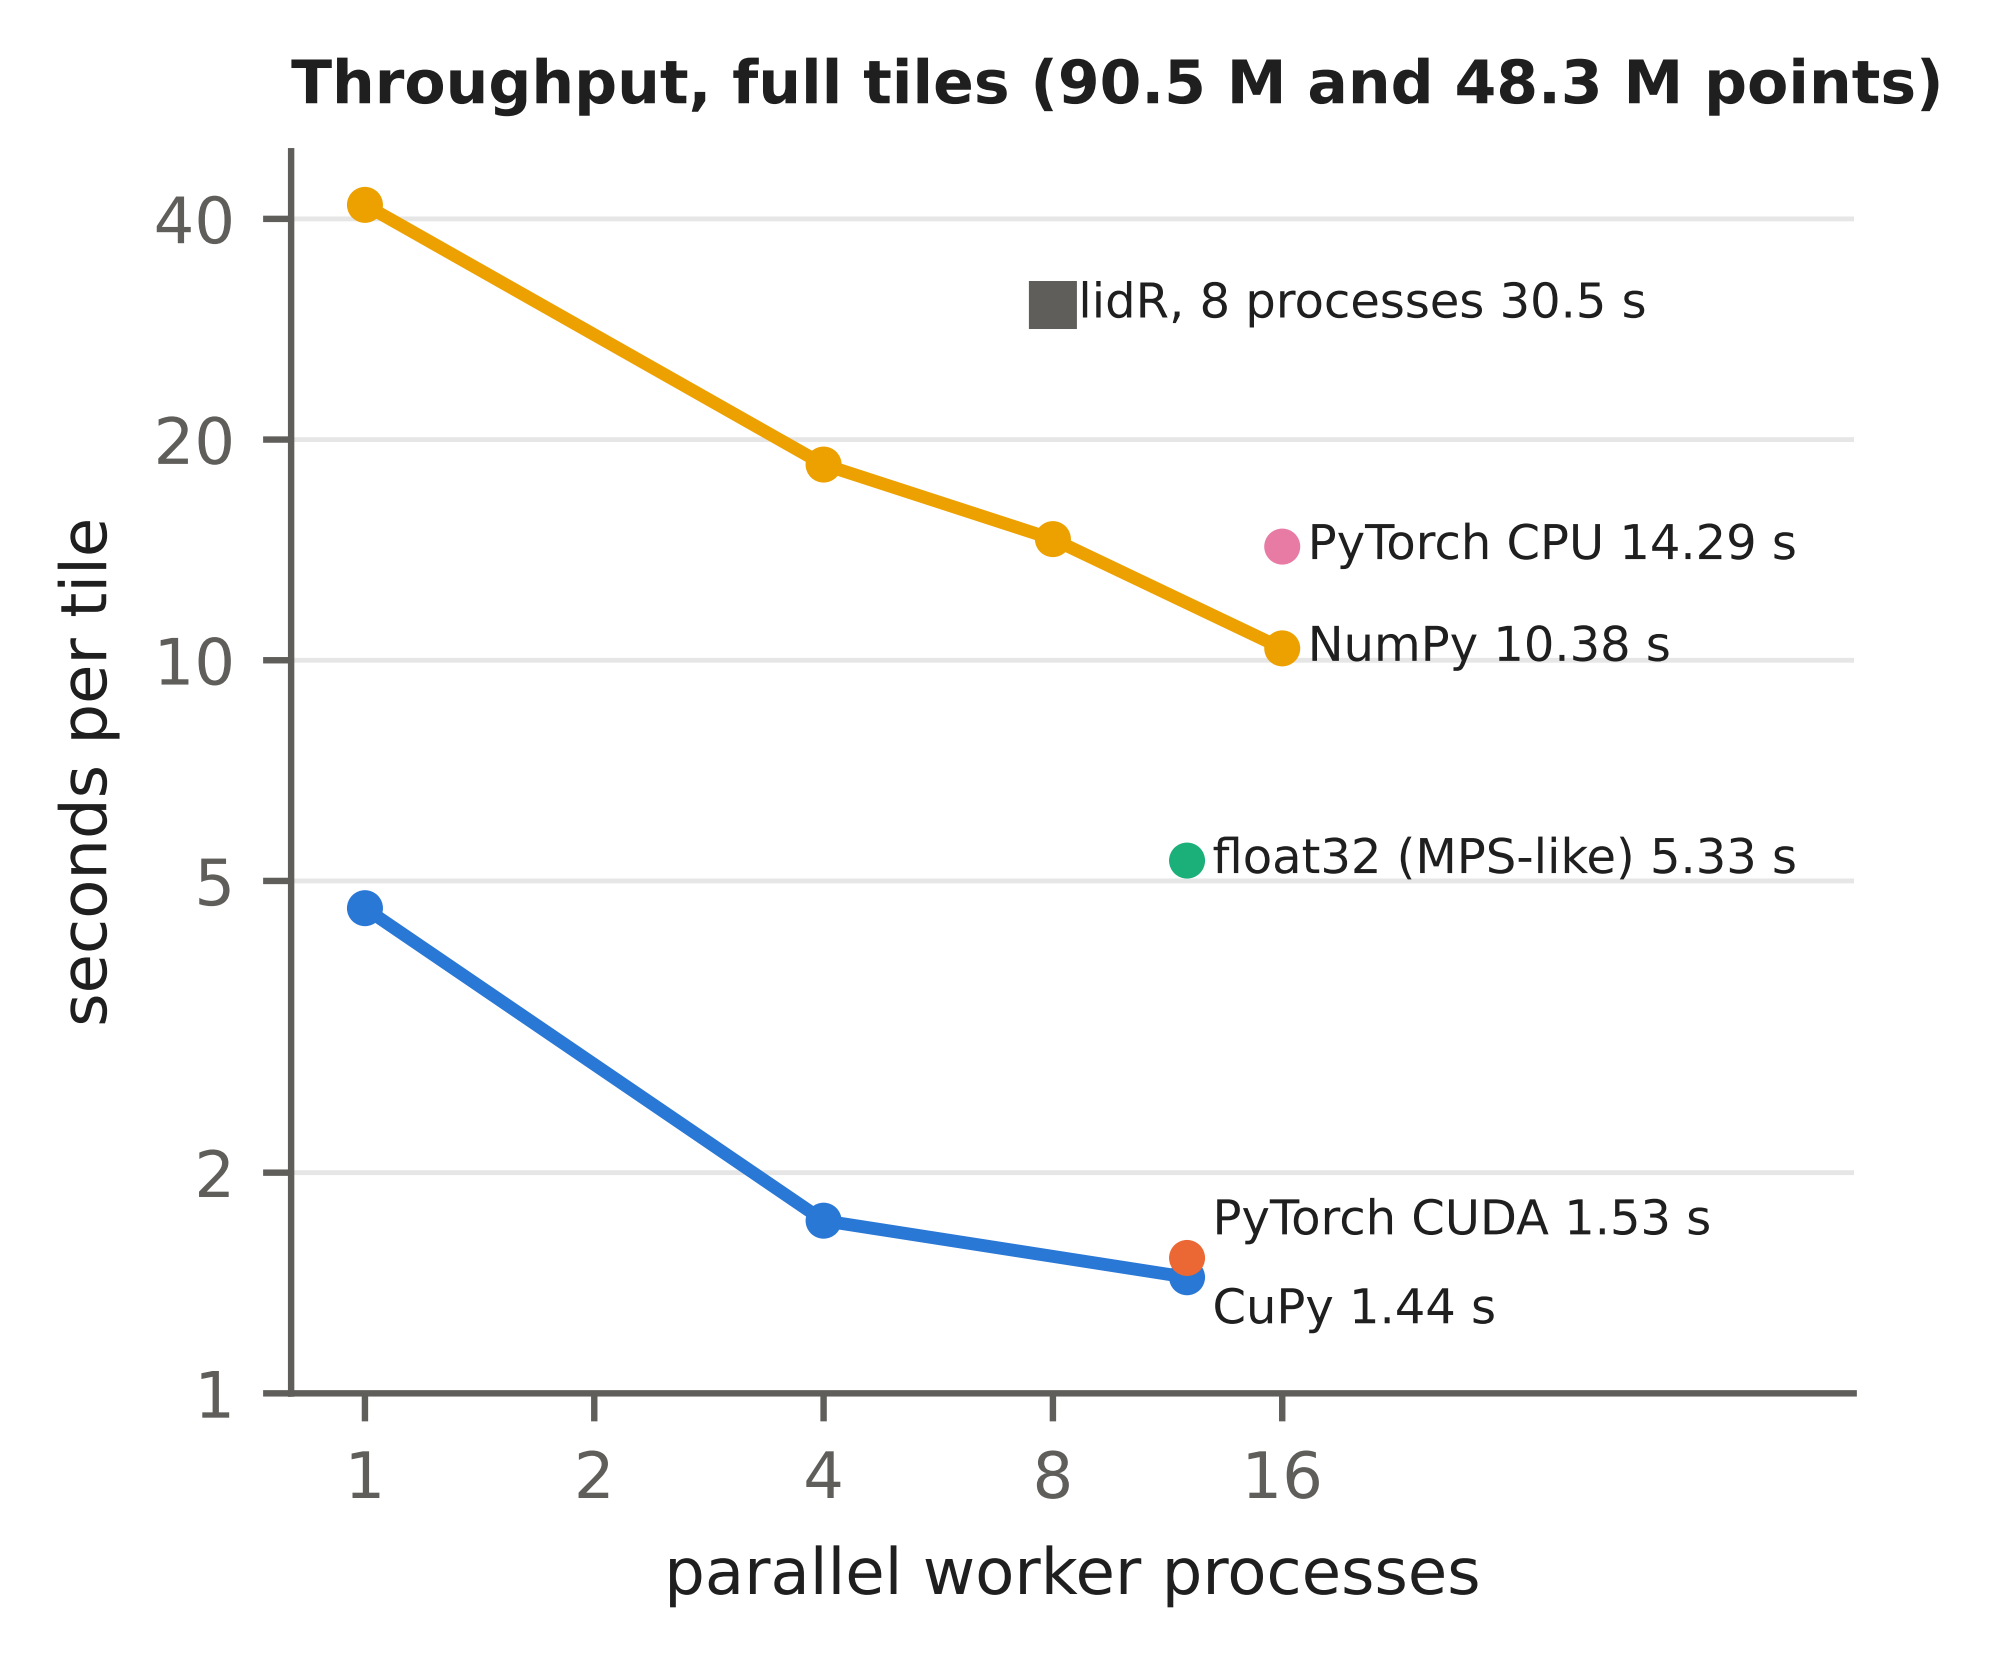
Throughput, full tiles (90.5 M and 48.3 M points)
lidR, 8 processes 30.5 s
PyTorch CPU 14.29 s
NumPy 10.38 s
float32 (MPS-like) 5.33 s
PyTorch CUDA 1.53 s
CuPy 1.44 s
seconds per tile
40
20
10
5
2
1
1
2
4
8
16
parallel worker processes